

ECE408/CS483/CSE408 Fall 2021

Applied Parallel Programming

Lecture 13:
Optimizing Convolutional Layers &
Final Project Kickoff

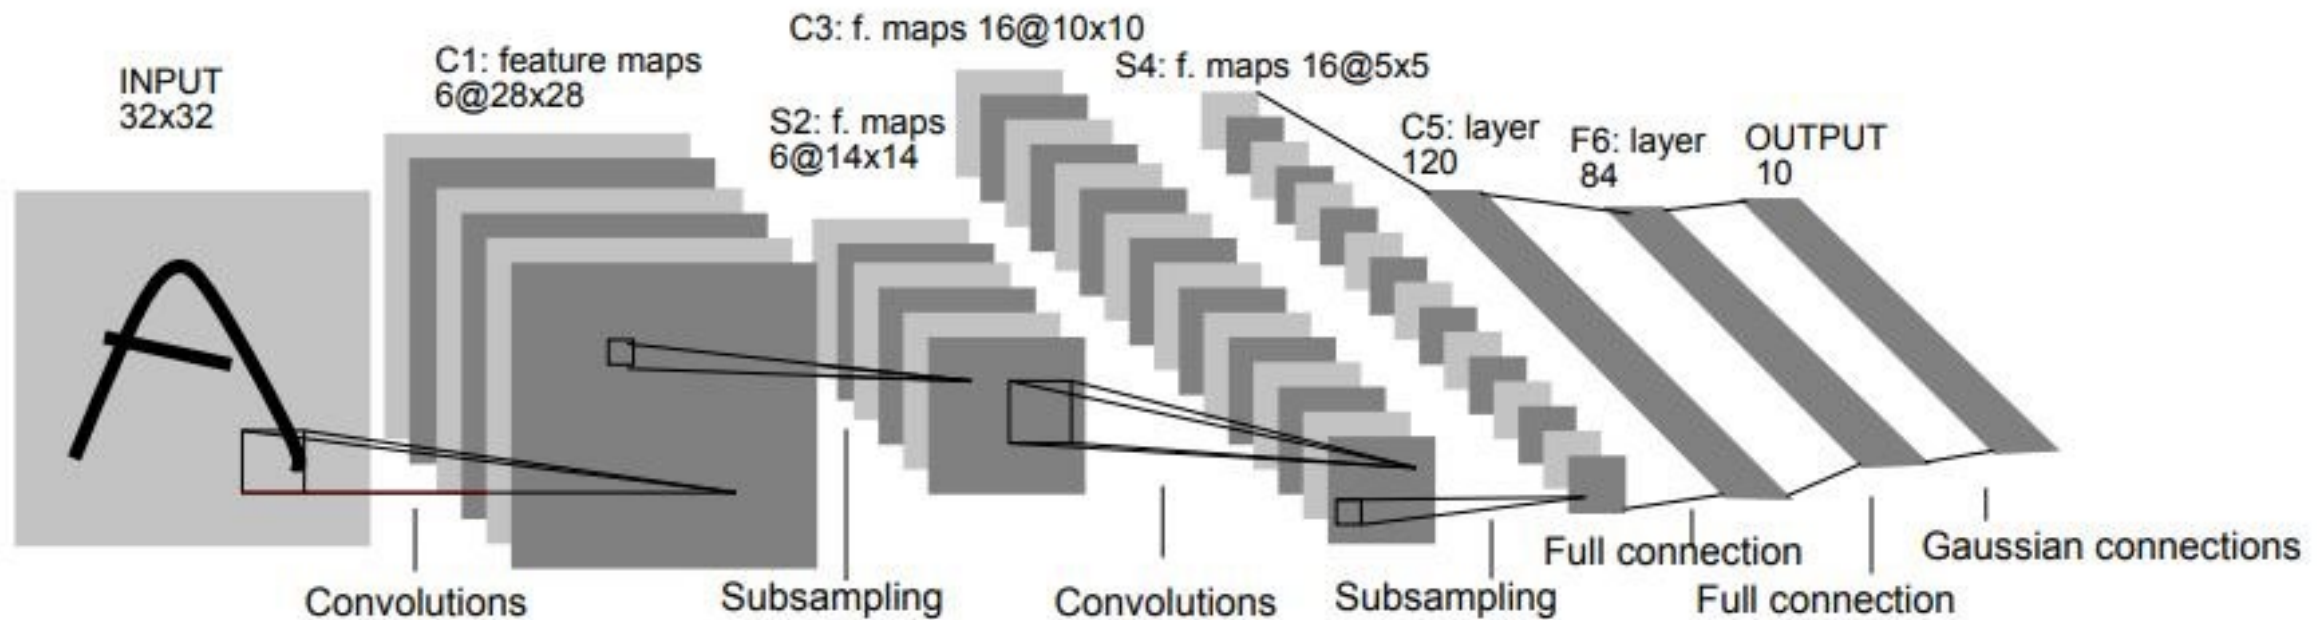
Course Reminders

- Midterm 1 is on this Thursday, October 7th
 - On-line, everybody will be taking it at the same time
 - Thursday, Oct. 7th 8:00pm-9:20pm US Central time
 - Friday, Oct. 9th 9:00am-10:20am Beijing time
 - Includes materials from Lecture 1 through Lecture 10
 - See <https://wiki.illinois.edu/wiki/display/ECE408/Exams> for details
- Project Milestone 1: Rai installation, CPU Convolution, Profiling
 - Due: Friday October 15th
 - Project details are posted on course wiki
<https://wiki.illinois.edu/wiki/display/ECE408/Project>
 - In today's lecture, we will go over these details

Objective

- To learn to implement the different types of layers in a Convolutional Neural Network (CNN) in CUDA
- To understand how unrolling input X can improve performance for convolution layers on GPUs
- Review Final Project

LeNet-5: CNN for hand-written digit recognition



Sequential Code: Forward Convolutional Layer

```
void cL_f(int B, int M, int C, int H, int W, int K, float* X, float* W, float* Y)
{
    int H_out = H - K + 1;           // calculate H_out, W_out
    int W_out = W - K + 1;

    for (int b = 0; b < B; ++b)      // for each image
        for(int m = 0; m < M; m++)   // for each output feature map
            for(int h = 0; h < H_out; h++) // for each output value (two loops)
                for(int w = 0; w < W_out; w++) {
                    Y[b, m, h, w] = 0.0f; // initialize sum to 0
                    for(int c = 0; c < C; c++) // sum over all input channels
                        for(int p = 0; p < K; p++) // KxK filter
                            for(int q = 0; q < K; q++)
                                Y[b, m, h, w] += X[b, c, h + p, w + q] * W[m, c, p, q];
                }
    }
```

Parallelism in a Convolution Layer

Output feature maps can be calculated in parallel

- Usually a small number, not sufficient to fully utilize a GPU

All **output** feature map **pixels** can be calculated in parallel

- All rows can be done in parallel
- All pixels in each row can be done in parallel
- Large number but diminishes as we go into deeper layers

All **input feature maps** can be processed in parallel,
but need atomic operation or tree reduction (we'll learn later)

Different layers may demand different strategies.

Sequential Code: Forward Pooling Layer

```
void pLayer_forward(int B, int M, int H_out, int W_out, int N, float* Y, float* S)
{
    for (int b = 0; b < B; ++b)                // for each image
        for (int m = 0; m < M; ++m)            // for each output feature map
            for (int x = 0; x < H_out/N; ++x)    // for each output value (two loops)
                for (int y = 0; y < W_out/N; ++y) {
                    float acc = 0.0f              // initialize sum to 0
                    for (int p = 0; p < N; ++p)    // loop over NxN block of Y (two loops)
                        for (int q = 0; q < N; ++q)
                            acc += Y[b, m, N*x + p, N*y + q];
                    acc /= N * N;                  // calculate average over block
                    S[b, m, x, y] = sigmoid(acc + bias[m]) // bias, non-linearity
                }
    }
```


Kernel Implementation of Subsampling Layer

- Straightforward mapping from grid to subsampled output feature map pixels
- in GPU kernel,
 - need to manipulate index mapping
 - for accessing the output feature map pixels
 - of the previous convolution layer.
- Often merged into the previous convolution layer to save memory bandwidth

Design of a Basic Kernel

- Each block computes
 - a tile of output pixels for one feature
 - `TILE_WIDTH` pixels in each dimension
- Grid's X dimension maps to M output feature maps
- Grid's Y dimension maps to the tiles in the output feature maps (linearized order).
- (Grid's Z dimension is used for images in batch, which we omit from slides.)

tiles covering an
output feature map,
marked with
linearized indices



0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19

A Small Example


Assume

- **M = 4** (4 output feature maps),
- thus 4 blocks in the X dimension, and
- **W_out = H_out = 8** (8x8 output features).

If **TILE_WIDTH = 4**,
we also need 4 blocks in the Y dimension:

- for each output feature,
- top two blocks in each column calculates the top row of tiles, and
- bottom two calculate the bottom row.

CUDA Grid and
Threadblocks



Output Feature Maps and Tiles 10

Host Code for a Basic Kernel: CUDA Grid

Consider an output feature map:

- width is **W_out**, and
- height is **H_out**.
- Assume these are multiples of **TILE_WIDTH**.

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19

Let **X_grid** be the number of blocks needed in X dim (5 above).

Let **Y_grid** be the number of blocks needed in Y dim (4 above).

Host Code for a Basic Kernel: CUDA Grid

(Assuming W_{out} and H_{out} are multiples of $TILE_WIDTH$.)

```
#define TILE_WIDTH 16          // We will use 4 for small examples.  
W_grid = W_out/TILE_WIDTH;    // number of horizontal tiles per output map  
H_grid = H_out/TILE_WIDTH;    // number of vertical tiles per output map  
Y = H_grid * W_grid;  
  
dim3 blockDim(TILE_WIDTH, TILE_WIDTH, 1); // output tile for untiled code  
dim3 gridDim(M, Y, 1);  
  
ConvLayerForward_Kernel<<< gridDim, blockDim >>>(...);
```

Partial Kernel Code for a Convolution Layer

```
__global__ void ConvLayerForward_Basic_Kernel
(int C, int W_grid, int K, float* X, float* W, float* Y)
{
    int m = blockIdx.x;
    int h = (blockIdx.y / W_grid) * TILE_WIDTH + threadIdx.y;
    int w = (blockIdx.y % W_grid) * TILE_WIDTH + threadIdx.x;
    float acc = 0.0f;
    for (int c = 0; c < C; c++) {                // sum over all input channels
        for (int p = 0; p < K; p++)                // loop over KxK filter
            for (int q = 0; q < K; q++)
                acc += X[c, h + p, w + q] * W[m, c, p, q];
    }
    Y[m, h, w] = acc;
}
```

Some Observations

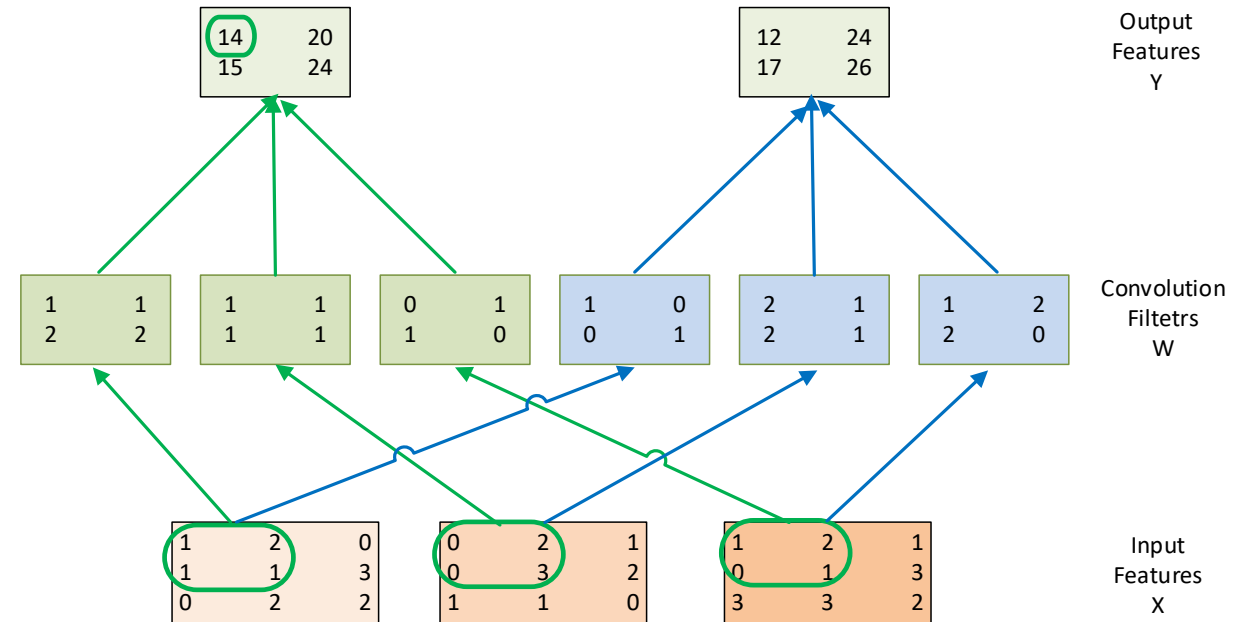
Enough parallelism

- if the total number of pixels
- across all output feature maps is large
- (often the case for CNN layers)

Each input tile

- loaded M times (number of output features), so
- **not efficient in global memory bandwidth,**
- but block scheduling in X dimension should give cache benefits.

Implementing a Convolution Layer with Matrix Multiplication



1	1	2	2
1	0	0	1

1	1	1	1
2	1	2	1

0	1	1	0
1	2	2	0

*

1	2	1	1
2	0	1	3
1	1	0	2
1	3	2	2
0	2	0	3
2	1	3	2
0	3	1	1
3	2	1	0
1	2	1	1
2	1	0	3
0	1	3	3
1	3	3	2

=


14	20	15	24
12	24	17	26

Convolution Filters W' Input Features X_{unrolled} Output Features Y

Simple Matrix Multiplication

Each product matrix element is an output feature map pixel.

This inner product generates element 0 of output feature map 0.

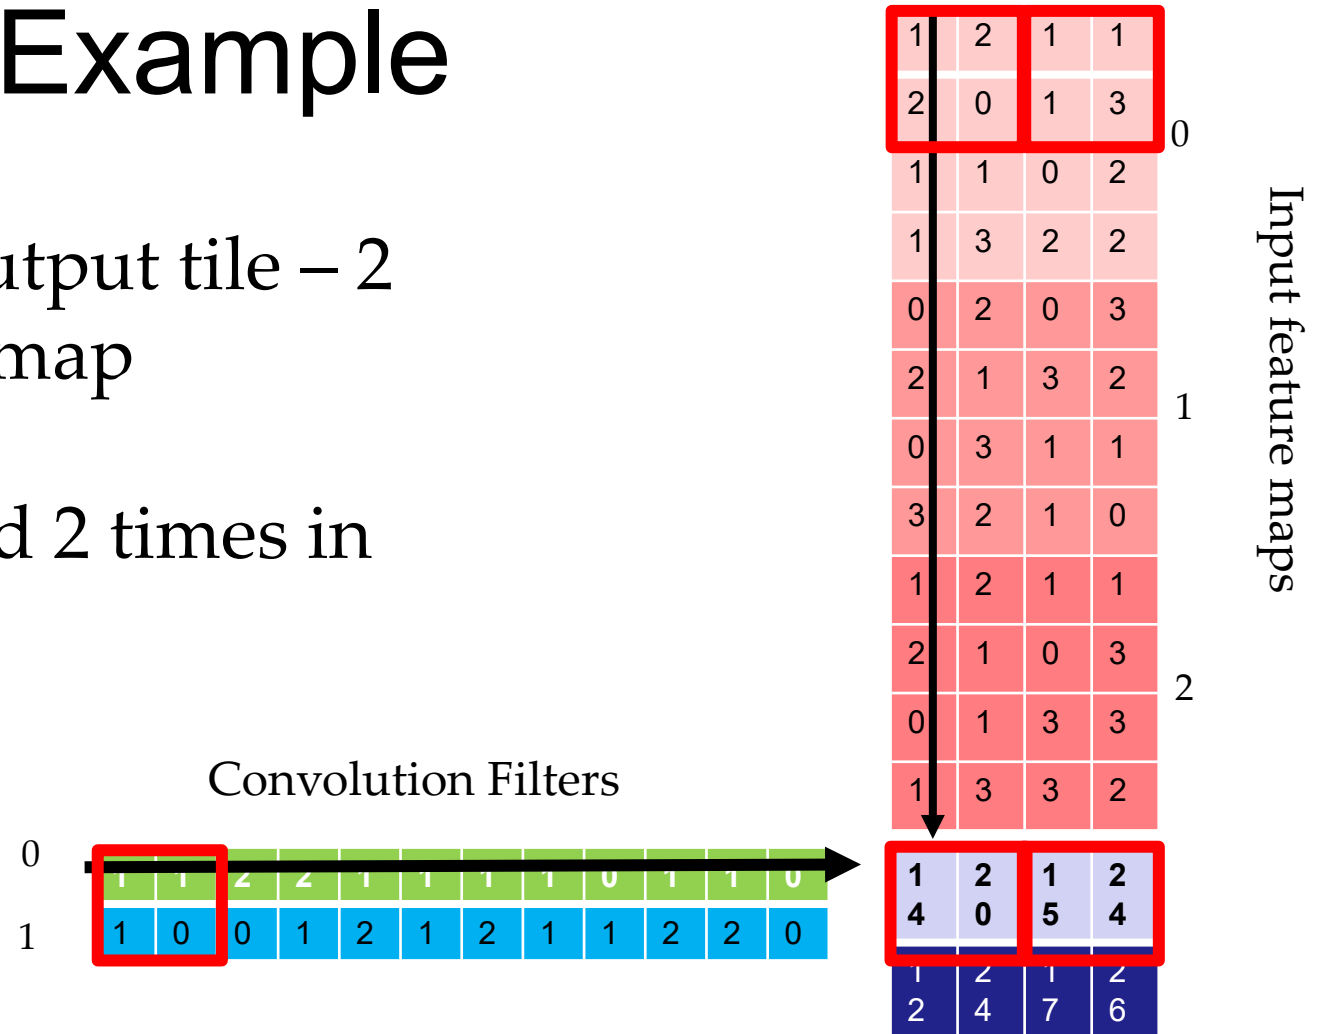


Tiled Matrix Multiplication

2x2 Example

Each block calculates one output tile – 2 elements from each output map

Each input element is reused 2 times in the shared memory

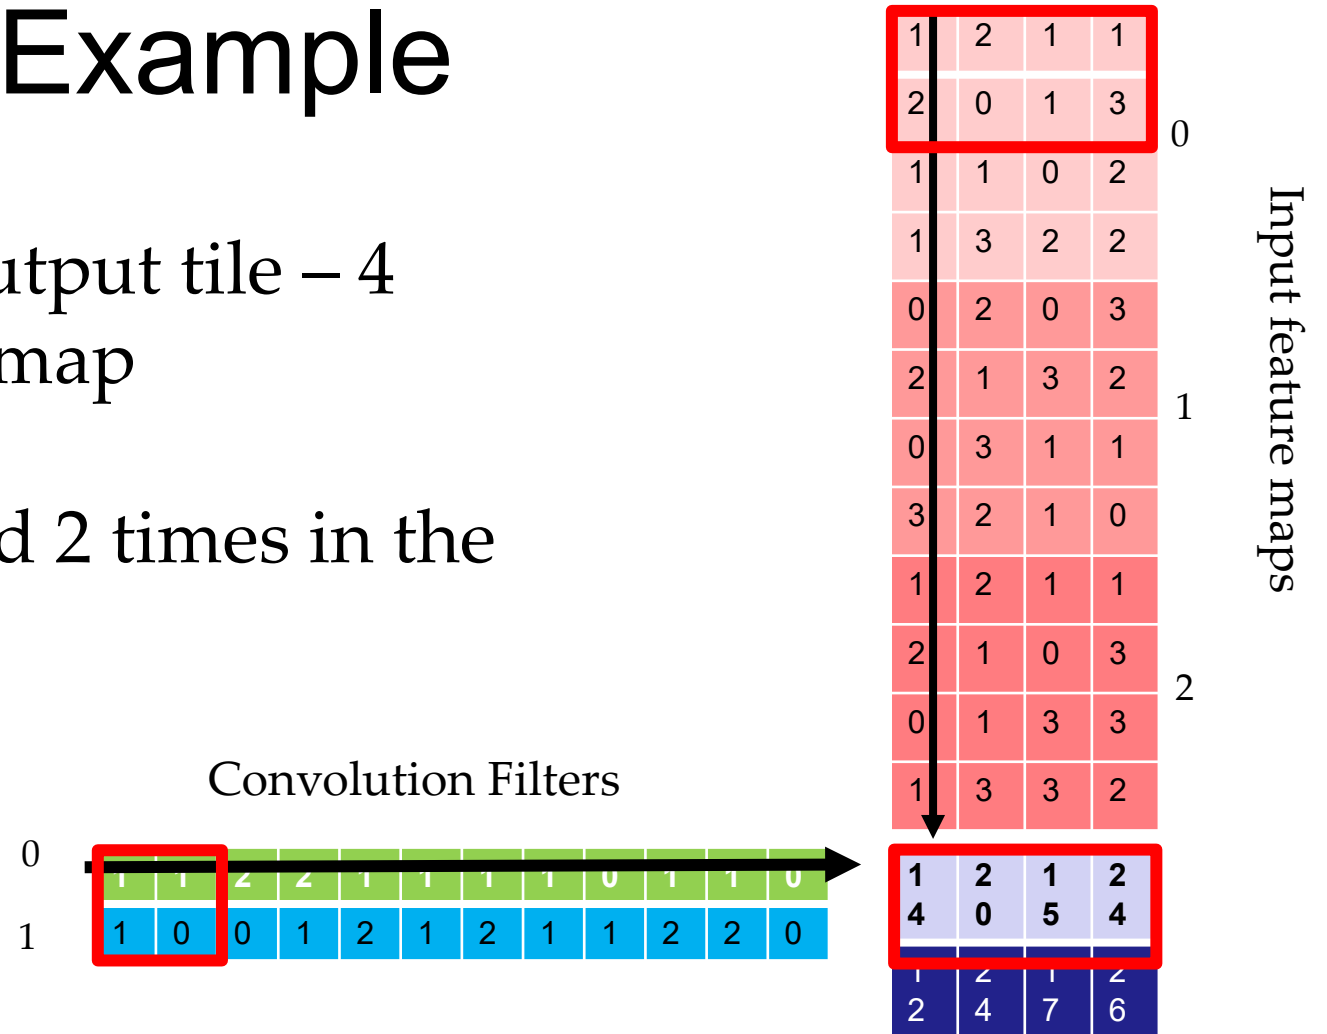


Tiled Matrix Multiplication

2x4 Example

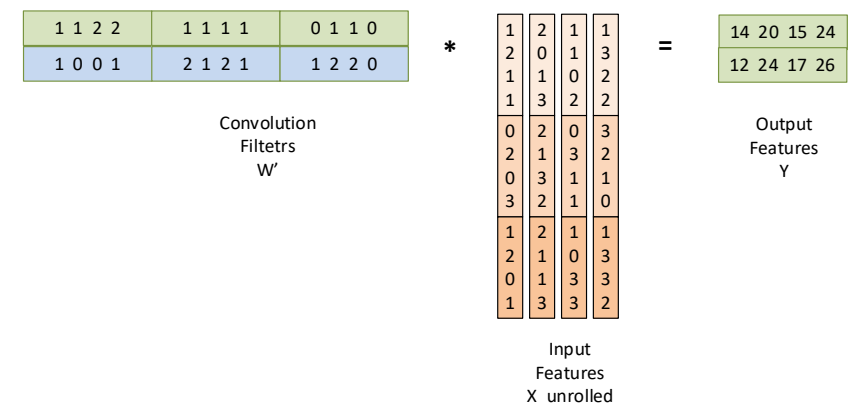
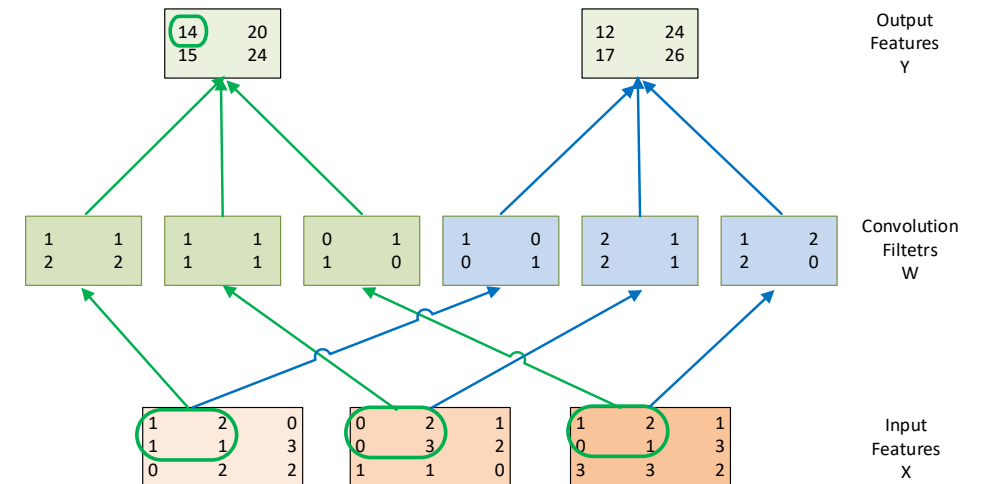
Each block calculates one output tile – 4 elements from each output map

Each input element is reused 2 times in the shared memory



Efficiency Analysis: Total Input Replication

- Replicated input features are shared among output maps
 - There are $H_{out} * W_{out}$ output feature map elements
 - Each requires $K*K$ elements from the input feature maps
 - So, the total number of input element after replication is $H_{out}*W_{out}*K*K$ times for each input feature map
 - The total number of elements in each original input feature map is $(H_{out}+K-1) * (W_{out}+K-1)$



Analysis of a Small Example

$$H_{\text{out}} = 2$$

$$W_{\text{out}} = 2$$

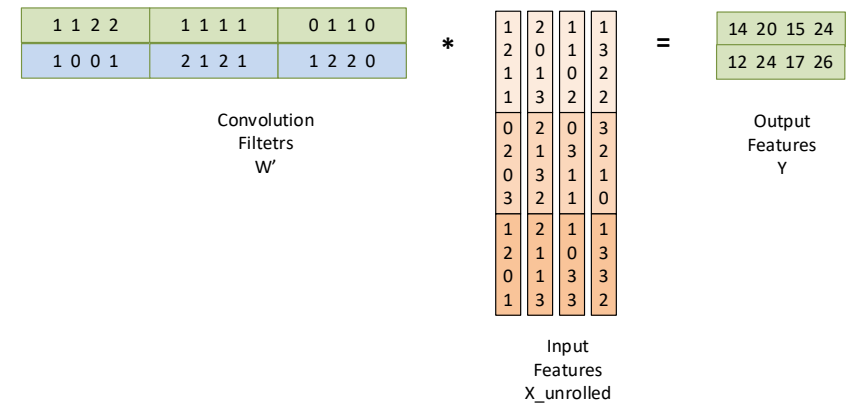
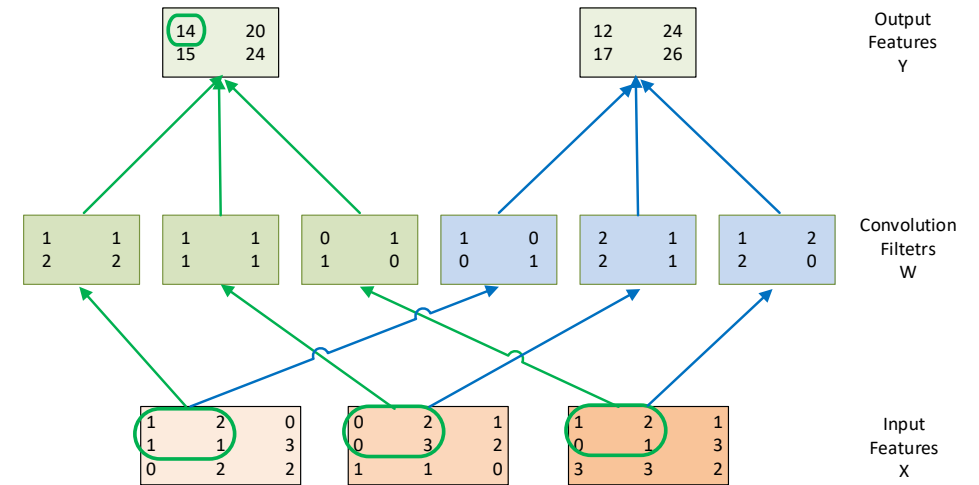
$$K = 2$$

There are 3 input maps (channels)

The total number of input elements in the replicated (“unrolled”) input matrix is $3 \times 2 \times 2 \times 2 \times 2$

The replicating factor is

$$(3 \times 2 \times 2 \times 2 \times 2) / (3 \times 3 \times 3) = 1.78$$

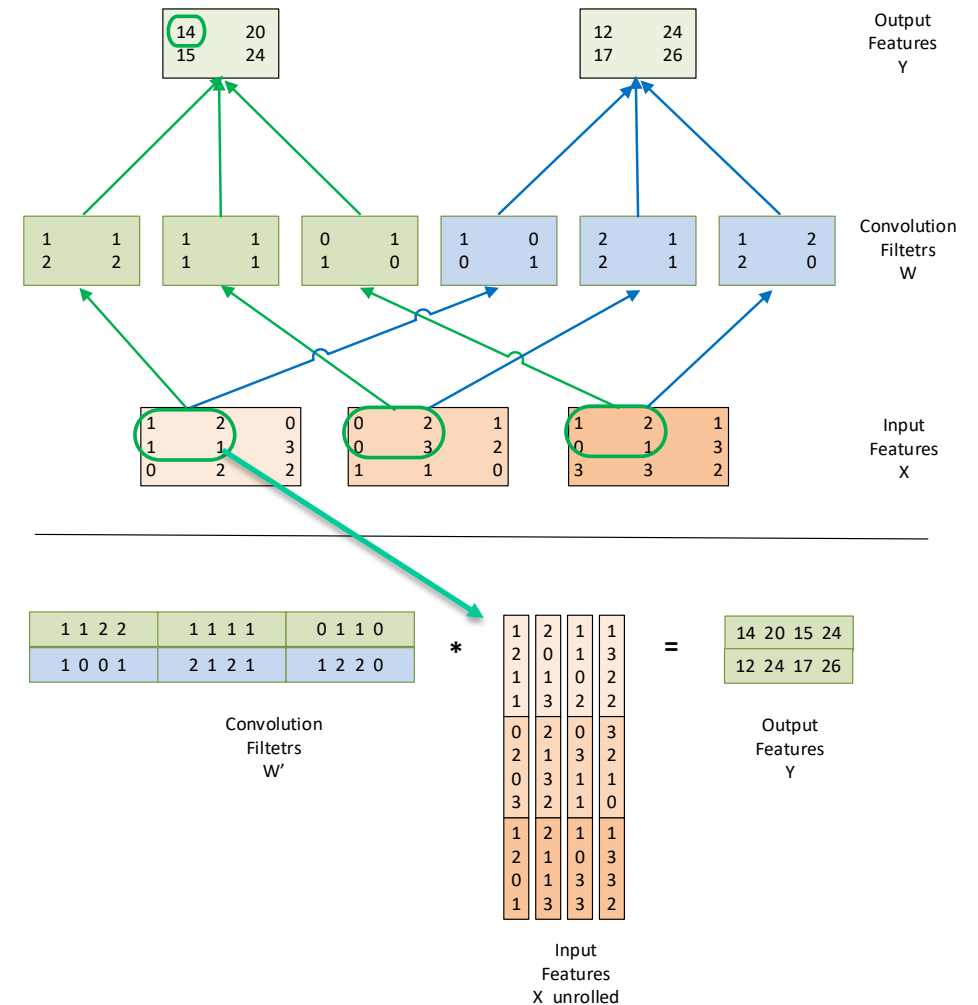


Memory Access Efficiency of Original Convolution Algorithm

- Assume that we use tiled 2D convolution
- For input elements
 - Each output tile has TILE_WIDTH^2 elements
 - Each input tile has $(\text{TILE_WIDTH}+K-1)^2$
 - The total number of input feature map element accesses was $\text{TILE_WIDTH}^2 * K^2$
 - The reduction factor of the tiled algorithm is $K^2 * \text{TILE_WIDTH}^2 / (\text{TILE_WIDTH}+K-1)^2$
- The convolution filter weight elements are reused within each output tile

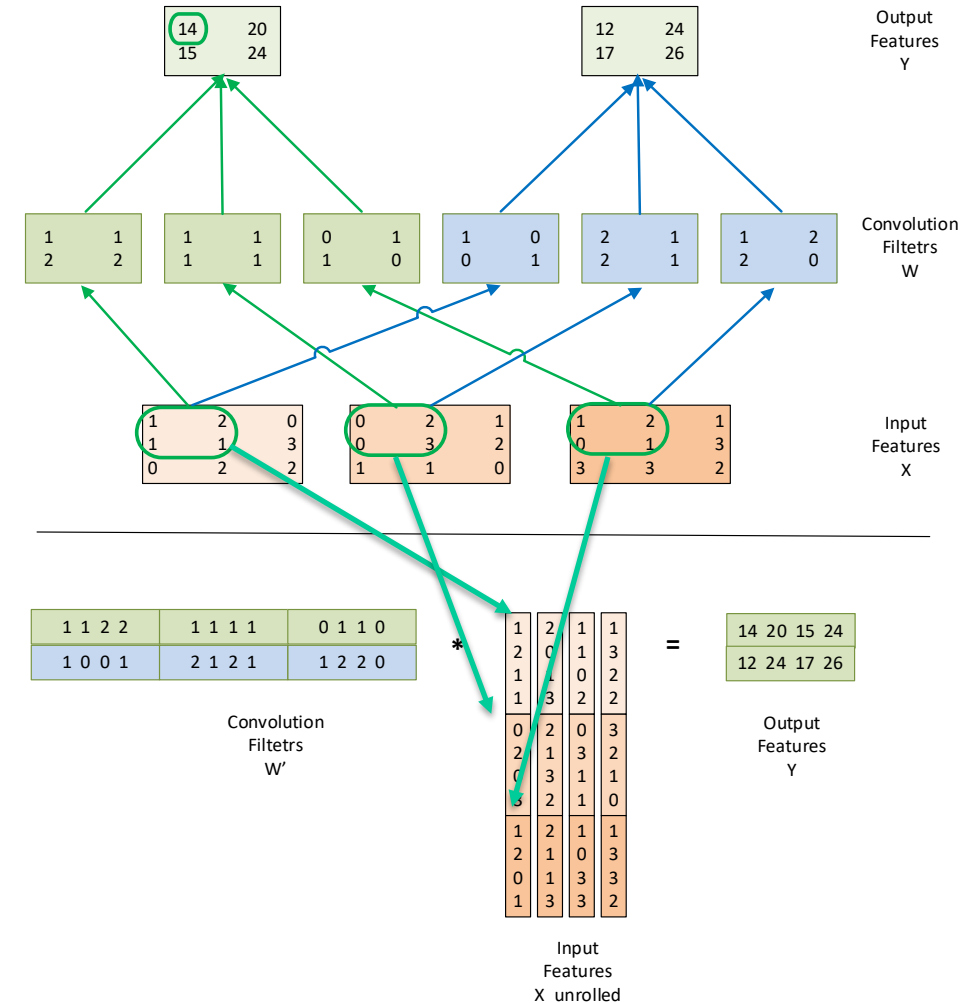
Properties of the Unrolled Matrix

- Each unrolled column corresponds to an output feature map element
- For an output feature element (h,w), the index for the unrolled column is $h * W_{out} + w$ (linearized index of the output feature map element)



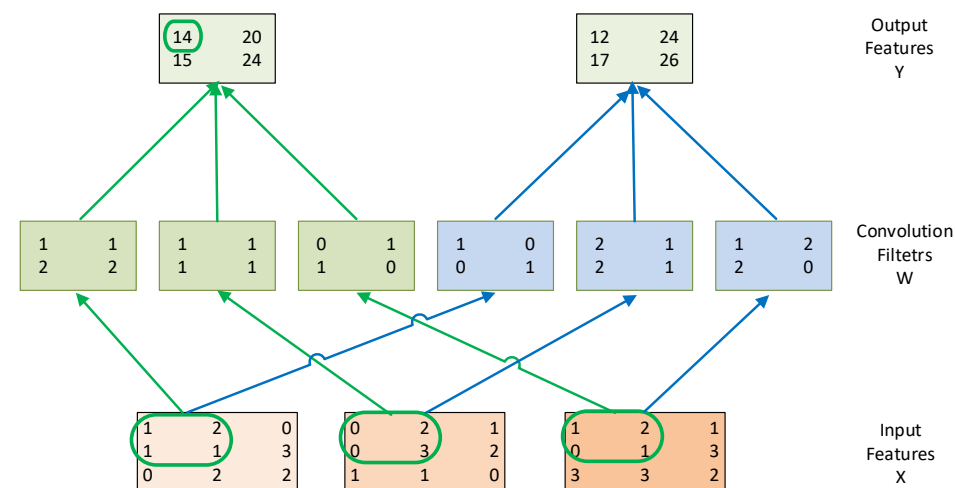
Properties of the Unrolled Matrix (cont.)

- Each section of the unrolled column corresponds to an input feature map
- Each section of the unrolled column has $k \times k$ elements (convolution mask size)
- For an input feature map c , the vertical index of its section in the unrolled column is $c \times k \times k$ (linearized index of the output feature map element)



To Find the Input Elements

- For output element (h,w) , the base index for the upper left corner of the input feature map c is (c, h, w)
- The input element index for multiplication with the convolution mask element (p, q) is $(c, h+p, w+q)$



1	1	2	2
1	0	0	1

1	1	1	1
2	1	2	1

0	1	1	0
1	2	2	0

*

1	2	1	1
2	0	1	3
1	1	0	2
1	3	2	2
0	2	0	3
2	1	3	2
0	3	1	1
3	2	1	0
1	2	1	1
2	1	0	3
0	1	3	3
1	3	3	2

=

14	20	15	24
12	24	17	26

Convolution Filters W'
Input Features $X_{unrolled}$
Output Features Y

Input to Unrolled Matrix Mapping

Output element (h, w)

Mask element (p, q)

Input feature map c

```
// calculate the horizontal matrix index
```

```
int w_unroll = h * W_out + w;
```

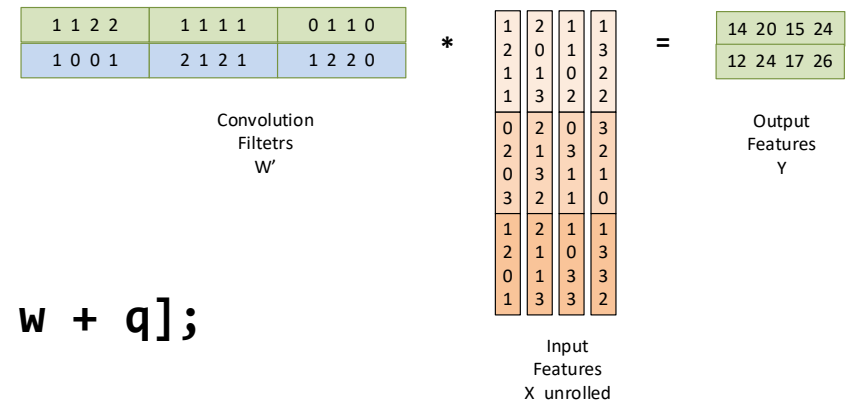
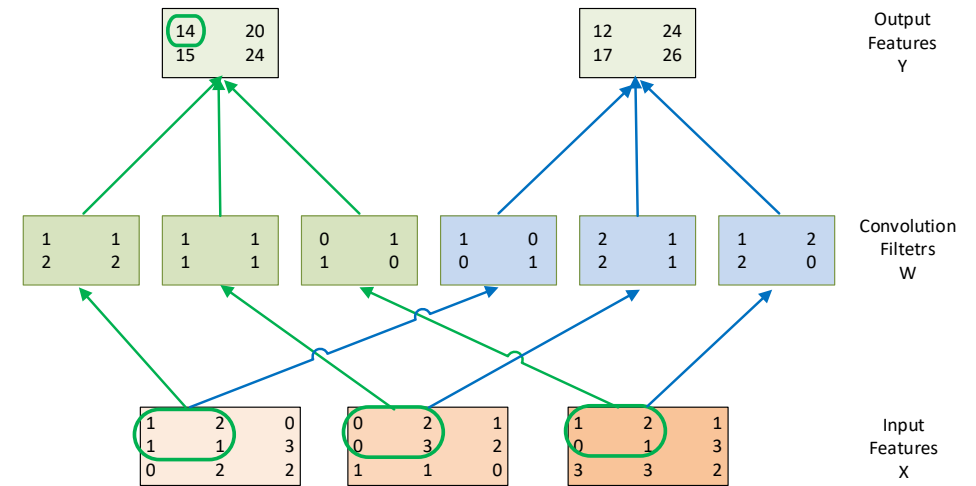
```
// find the beginning of the unrolled
```

```
int w_base = c * (K*K);
```

```
// calculate the vertical matrix index
```

```
int h_unroll = w_base + p * K + q;
```

```
X_unroll[b, h_unroll, w_unroll] = X[b, c, h + p, w + q];
```



Function to generate “unrolled” X

```
void unroll(int B, int C, int H, int W, int K, float* X, float* X_unroll)
{
    int H_out = H - K + 1;           // calculate H_out, W_out
    int W_out = W - K + 1;
    for (int b = 0; b < B; ++b)      // for each image
        for (int c = 0; c < C; ++c) { // for each input channel
            int w_base = c * (K*K);  // per-channel offset for smallest X_unroll index
            for (int p = 0; p < K; ++p) // for each element of KxK filter (two loops)
                for (int q = 0; q < K; ++q) {
                    for (int h = 0; h < H_out; ++h) // for each thread (each output value, two loops)
                        for (int w = 0; w < W_out; ++w) {
                            int h_unroll = w_base + p * K + q; // data needed by one thread
                            int w_unroll = h * W_out + w;       // smallest index--across threads (output values)
                            X_unroll[b, h_unroll, w_unroll] = X[b, c, h + p, w + q]; // copy input pixels
                        }
                }
        }
}
```

Implementation Strategies for a Convolution Layer

- **Baseline**
 - Tiled 2D convolution implementation, use constant memory for convolution masks
- **Matrix-Multiplication Baseline**
 - Input feature map unrolling kernel, constant memory for convolution masks as an optimization
 - Tiled matrix multiplication kernel
- **Matrix-Multiplication with built-in unrolling**
 - Perform unrolling only when loading a tile for matrix multiplication
 - The unrolled matrix is only conceptual
 - When loading a tile element of the conceptual unrolled matrix into the shared memory, use the properties in the lecture to load from the input feature map
- **More advanced Matrix-Multiplication**
 - Use joint register-shared memory tiling

Two vertical lines, one blue and one orange, are positioned on the left side of the slide.

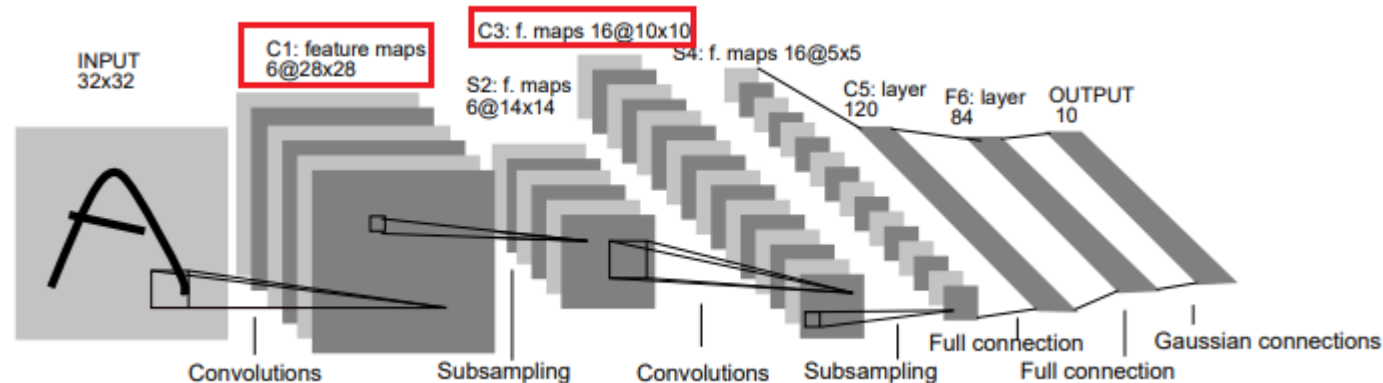
**ANY MORE QUESTIONS?
READ CHAPTER 16**

Project Overview

- Optimize the forward pass of the convolutional layers in a modified LeNet-5 CNN using CUDA. (CNN implemented using Mini-DNN, a C++ framework)
- The network will be classifying Fashion MNIST dataset
- Some network parameters to be aware of
 - Input Size: 86x86 pixels, batch of 10k images
 - Input Channels: 1
 - Convolutional kernel size: 7x7
 - Number of kernels: Variable (your code should support this)



<https://github.com/zalandoresearch/fashion-mnist>



Project Timeline

- **All milestones are due on Fridays at 8 pm Central Time**
- Everyone must individually submit all Milestones.
 - **No sharing of code is allowed**
- October 15th: Project milestone 1:
 - Rai installation, CPU Convolution, profiling
- November 5th: Project milestone 2:
 - Baseline GPU Convolution Kernel
- December 3rd: Project milestone 3:
 - GPU Convolution Kernel Optimizations


Project Release

- Project is released now (only PM1 for now)
 - Check the course wiki page for the link to the github repository
- Project Landing Page:
 - <https://wiki.illinois.edu/wiki/display/ECE408/Project>
- Make sure that you clone the **2021fa** branch using git

```
git clone -b 2021fa https://github.com/illinois-impact/ece408_project.git
```
- The readme in the repository contains all the instructions and details to complete the project.

The RAI Submission System

- RAI is a scalable job submission system designed for CPU and GPU workloads



Setting up the RAI Client

- The link to download the client is in the project readme. Download the client version based on your OS.
- You will receive a `.rai_profile` file via email
 - Place this file in your home directory (i.e., `~/rai_profile` on Linux/macOS)
 - This file will contain the credentials required for you to submit jobs via RAI

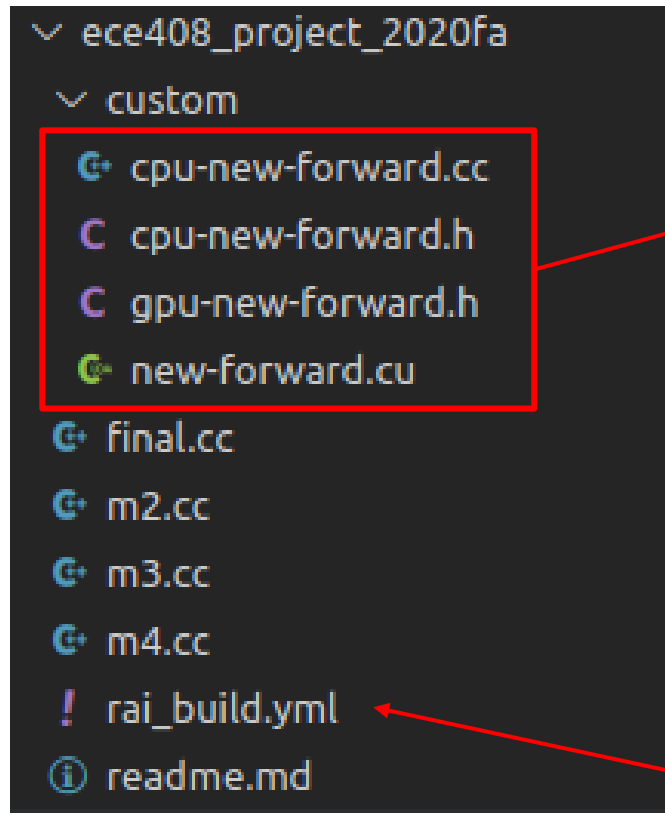
```
profile:
  id: 5f737b0788a5ec178e188f64
  firstname: James
  lastname: Cyriac
  username: jcyriac2
  email: jcyriac2@illinois.edu
  access_key: <snip>
  secret_key: <snip>
  affiliation: uiuc
```

Scheduling a job using RAI

```
rai -p <project folder> --queue rai_amd64_ece408
```

- Make sure the project folder points to the root of the Github repository
- Uploads your folder to the file server and queues your work request
- The code in `rai_build.yml` is executed on the server in the specified docker container
- The results are streamed back to you in real time

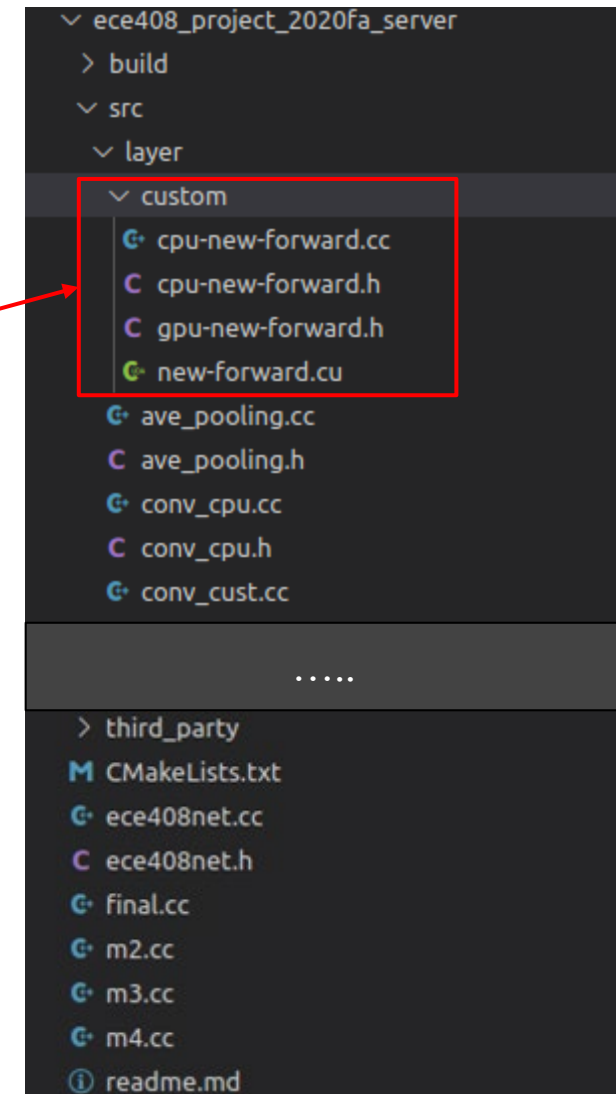
Mini-DNN Code Structure



Client Side

RAI replaces these 4 files on the server, recompiles the CNN and runs the inference

The commands RAI will run is specified in `rai_build.yml`. You will need to modify this file.



Server Side

Submitting your code

To submit your code for grading, run

```
rai -p <project folder> --queue rai_amd64_ece408  
--submit=[m1, m2, m3]
```

We will be using a specific version of `rai_build.yml` when grading. (We will not use the `rai_build.yml` in your project directory)

Make sure to include your `report.pdf` for the respective milestone in the project folder when you submit.

RAI tips and tricks

- `rai -p <project folder> --queue rai_amd64_ece408 history`

Shows your last 20 RAI runs and the associated file URL where your build directory is stored

- `rai -p <project folder> --queue rai_amd64_ece408 queued`

Shows the number of jobs in the queue and are waiting to start

- `rai -p <project folder> --queue rai_amd64_ece408 ranking`

Shows anonymized performance results comparing your team to other teams.
Only applicable from Milestone 3

- `rai --help` - Lists all the other commands that can be run using RAI